

## **Lecture 7**

- Demonstration of the resources available on Prince HPC cluster
- Fourier methods for financial models and applications

## 7.1 Accessing Prince HPC Cluster and Using CUDA C on Linux Operating System

- Go through steps for setting up and logging into HPC cluster

<https://wikis.nyu.edu/display/NYUHPC/High+Performance+Computing+at+NYU>

<https://wikis.nyu.edu/display/NYUHPC/Clusters+-+Prince>

- Do Nsight Eclipse IDE, and show code examples
  - multi-thread in C++ on Linux is a little different: use pthread
- Run examples on Prince HPC: multi-threading example that uses multiple GPU's
- Executing batch jobs on Prince

## 7.1 Prince Cluster and CUDA C on Linux Operating System (cont.)

- Logging into Prince cluster is a 2 step process that requires your NYU netid. There is information on the NYU network - see web pages
- First, you need to download Putty, Xming, and WinSCP, if you plan to log in to Prince from a Windows machine
- Open Xming by clicking on icon. It runs in the background

First, you need to open a tunnel using Putty. Instructions on how to set this up are on web site. Leave this screen in the background. Then you start a 2<sup>nd</sup> Putty. Again, the web site has instruction for setting this up. This is a somewhat complex 2 step process, but it works. This places in a Linux terminal session on the HPC cluster.

- → Examples

## 7.1 Prince Cluster and CUDA C (cont.)

- Use WinSCP to transfer files to your space on the Prince. On left panel, click on “Transferring data to/from Prince cluster

got to <https://wikis.nyu.edu/pages/viewpage.action?pageId=84612833>

- You will need to use Linux commands on the cluster

ls	list directories and files in current directory
ls -la	list directories and files, with file sizes
cd <b>name</b>	change to directory, may need full path, cd ~/MyTestFiles/SourceCode
cd ..	Go back to the next highest directory
mkdir <b>name</b>	make a new directory, in current location
rm <b>filename</b>	removes, or deletes <b>filename</b> (be careful)
rmdir <b>name</b>	removes directory <b>name</b>
cat <b>filename</b>	prints <b>filename</b> to the screen
less <b>filename</b>	prints <b>filename</b> to the screen, in pages (space bar to continue, q to quit)
cp	copy files, cp <b>filename1</b> ~/Directory/ <b>filename2</b>

You can use man command to retrieve information on Linux commands

Highlight text and copy with right-click if you are on Windows

## 7.1 Prince Cluster and CUDA C (cont.)

- To build or make executable files to run C++ or Cuda C programs, use Nsight Eclipse. You can copy C++ and Cuda C programs to Prince cluster, but you must do some editing and rebuild on Prince (Windows executables do not run on Linux systems)
- Copy your source code and parameter text files to Prince using WinSCP. I recommend editing text files with gedit. Use Nsight to edit and rebuild C++ and Cuda C programs. Linux systems also have Emacs
- Use the following commands to set up cuda on Prince.

```
$ module avail    → to see modules available  
$ module load cuda/10.0.130
```

- Version 8 and 9 are also available. Use `$ module spider cuda` to get more information.
- Start the Nsight Eclipse IDE by executing the following command

```
$ nsight
```

- This GUI can be slow to start up.

## 7.1 Prince Cluster and CUDA C (cont.)

- Open Nsight
- To start a new C++ project, right click the tab at the top left of Nsight and choose New → C++ project
  - 1) on the next screen, choose Empty Project and Linux GCC,, and type in a project name
  - 2) in the directory created for this project, create an src directory ([mkdir src](#))
  - 3) copy your source code and header files into the src directory
  - 4) open the project
  - 5) recommend: right-click on the project and click on “Close unrelated projects”
  - 6) right click on project and go to Properties at the bottom of menu. From the next screen, choose Build → Settings. For multi-threading, you need to add -pthread to G++ and G++ Linker.
  - 7) After completing edits, use Save All. Do a Build using Debug and do a second Build using Release. You will need to add -pthread to the Build Settings again (both G++ and G++ Linker) for release
- To start a new Cuda C project, choose New → Cuda C/C++ project
  - 1) on the next screen, choose Empty Project and Cuda C, and type in a project name
  - 2) on the next screen, check 3.7 for the K80 and 7.0 for the V100 in both rows
  - 3) in the directory created for this project, create an src directory ([mkdir src](#))
  - 4) copy your source code and header files into the src directory
  - 5) open the project
  - 6) recommend: right-click on the project and click on “Close unrelated projects”
  - 7) After completing edits, use Save All. Do a Build using Debug and do a second Build using Release.

## 7.1 Prince Cluster and CUDA C (cont.)

### *Standard Compute Nodes*

- 4 nodes (Dell PowerEdge C6420 in a 6400 chassis enclosure) each with 2 Intel Xeon Gold 6148 2.4GHz CPUs ("Skylake", 20 cores/socket, 40 cores/node) and **187GB** memory, EDR interconnects. Nodes: **c42-0[1-4]**
- 68 nodes each with 2 Intel Xeon E5-2690v4 2.6GHz CPUs ("Broadwell", 14 cores/socket, 28 cores/node) and **125GB** memory, EDR interconnects
- 32 nodes each with 2 Intel Xeon E5-2690v4 2.6GHz CPUs ("Broadwell", 14 cores/socket, 28 cores/node) and **250GB** memory, EDR interconnects
- 32 nodes each with 2 Intel Xeon E5-2660v3 2.6GHz CPUs ("Haswell", 10 cores/socket, 20 cores/node) and **62 GB** memory. The 32 nodes are M630 Blade servers on 2 M1000e chassis and are interconnected via FDR Infiniband
- 64 nodes each with 2 Intel Xeon E5-2690v2 3.0GHz CPUs ("Ivy Bridge", 10 cores/socket, 20 cores/node) and **62GB** memory. The 64 nodes are M620 Blade servers on 4 M1000e chassis and are interconnected via FDR Infiniband (used to be Mercer chassis 0, 1, 2, 3)
- 112 nodes each with 2 Intel Xeon E-2690v2 3.0GHz CPUs ("Ivy Bridge", 10 cores/socket, 20 cores/node) and **62GB** memory. The 112 nodes are M620 Blade servers on 7 M1000e chassis and are interconnected via QDR Infiniband (Mercer chassis 14-20)
- 48 nodes each with 2 Intel Xeon E-2690v2 3.0GHz CPUs ("Ivy Bridge", 10 cores/socket, 20 cores/node) and **189GB** memory. The 48 nodes are M620 Blade servers on 3 M1000e chassis and are interconnected via QDR Infiniband (Mercer chassis 21-23)

### *Nodes equipped with NVIDIA GPUs*

- 6 nodes each with 2 Intel Xeon Gold 6148 2.4GHz CPUs ("Skylake", 20 cores/socket, 40 cores/node) and 384GB memory, EDR interconnects, each node equipped with **4 NVIDIA V100 SXM2 GPUs (16GB)** connected with NVLink
- 1 node with 2 Intel Xeon Gold 6148 2.4GHz CPUs ("Skylake", 20 cores/socket, 40 cores/node) and 192GB memory, EDR interconnects, each node equipped with **2 NVIDIA V100 PCIe GPUs (16GB)** connected via PCIe
- 8 nodes each with 2 Intel Xeon E5-2690v4 2.6GHz CPUs ("Broadwell", 14 cores/socket, 28 cores/node) and 256GB memory, EDR interconnects, each node equipped with **4 NVIDIA P100 GPUs (16GB)**
- 24 nodes each with 2 Intel Xeon E5-2690v4 2.6GHz CPUs ("Broadwell", 14 cores/socket, 28 cores/node) and 256GB memory, EDR interconnects, each node equipped with **4 NVIDIA P40 GPUs (24GB)**
- 9 nodes each with 2 Intel Xeon E5-2690v4 2.6GHz CPUs ("Broadwell", 14 cores/socket, 28 cores/node) and 256GB memory, EDR interconnects, each node equipped with **2 NVIDIA K80 GPUs (24GB, split between 2 GPU cards)**
- 8 nodes each with 2 Intel Xeon E5-2670v2 2.5GHz CPUs ("Ivy Bridge", 10 cores/socket, 20 cores/node) and 128 GB memory, FDR interconnects, each node equipped with **4 NVIDIA K80 GPUs**
- 4 nodes each with 2 Intel Xeon E5-2690v4 2.6GHz CPUs ("Broadwell", 14 cores/socket, 28 cores/node) and 128GB memory, EDR interconnects, each node equipped with **4 NVIDIA GTX 1080 GPUs (8 GB)**

## 7.1 Prince Cluster and CUDA C (cont.)

- Go to class examples on Prince cluster
- Recommended optimizations for g++ or gcc, for code on CPU only (on Linux machines)  
-O3 -ffast-math -march=native -funroll-loops
- Recommended optimizations for Cuda C, for code on GPU  
  
use -O3 for CPU or host code  
choose -O3 and fast math option
- Optimizations in Windows Visual Studio
  - For C++ CPU only code, go to Properties or Property Pages, and choose C/C++ → Optimization → Maximize speed (/O2)
  - For programs running on GPU, under Cuda C, choose Yes for Use Fast Math and Maximize Speed (/O2) for Optimization



## 7.1 Prince Cluster and CUDA C (cont.)

- One CANNOT run GPU executables on Prince. The Prince cluster is set up for batch jobs only.
- See the web site for instructions on how to set up batch jobs using \*.sh files

<https://wikis.nyu.edu/display/NYUHPC/Submitting+jobs+with+SBATCH>

- From the command line you type `$ sbatch filename.sh`
- For information on Prince GPU's go to

<https://wikis.nyu.edu/display/NYUHPC/Clusters+-+Prince>

- To access GPU's , add the following lines to batch file  
    `#SBATCH --gres=gpu:1` or `#SBATCH --gres=gpu:v100:1` to use a V100  
    `#SBATCH --gres=gpu:k80:2` will set up batch job to use 2 GPU's
- Do examples of batch files, and run one on Prince

## 7.1 Prince Cluster and CUDA C (cont.)

- LOSBatch\_GPU.sh -- batch file to run a multi-threaded program on 2 GPU's

```
#!/bin/bash
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=3
#SBATCH --time=5:00:00
#SBATCH --mem=4GB
#SBATCH --job-name=LOS_Test_HW_Model_GPU
#SBATCH --mail-type=END
#SBATCH --mail-user=ls127@nyu.edu
#SBATCH --output=slurm_%j.out
```

```
module purge
cd ~/MytestRuns
Release/Test_FD_LIBOROption_GPU
```

```
RUNDIR=$SCRATCH/my_project/run-`${SLURM_JOB_ID/.}*}
mkdir -p $RUNDIR
```